

Globalised Learned Pruning: Weight-Conditioned Hypernetwork for Structured Neuron Pruning of Frozen Language Models

Keshava Prasad

September 2, 2026

Abstract

Structured pruning of large language models is usually driven by hand-designed, data-dependent scoring rules: weight or activation magnitude, first-order Taylor estimates, or one-shot calibration against a closed-form reconstruction objective. These methods score neurons independently and layer-by-layer, with no mechanism for one layer’s redundancy to inform another’s. We propose **Globalised Learned Pruning (GLP)**: a small hypernetwork: a per-neuron weight-row encoder plus a bidirectional LSTM that shares context across network depth, that reads a frozen, pretrained model’s own weight matrices and emits binary keep/prune gates for feed-forward neurons, trained end-to-end against a soft sparsity penalty λ on the frozen model’s own next-token loss. The base model is never fine-tuned, before or after pruning; the deliverable is a smaller *effective* model at a controlled, characterized accuracy cost, not an improved one.

A toy-scale pilot (MLPs on MNIST/CIFAR-10) establishes, qualitatively, that weight matrices alone, no data, no gradients, and no retraining, are sufficient to find a highly compressible subnetwork, that the discovered mask is a near-deterministic property of the trained weights, and that prunability tracks layer width rather than depth. At LLM scale, the same architecture and hyperparameters, unmodified beyond the weight-extraction plumbing, port across GPT-2 small, OPT-125M, Mistral-7B, and Llama-2-7B. On GPT-2 and OPT-125M we report clean sparsity-perplexity Pareto frontiers and show the learned pruner cuts perplexity cost by roughly $6\times$ versus a matched-sparsity activation-magnitude baseline. At 7B scale we validate directly against a recently published trained-hypernetwork baseline, DISP-LLM (Gao et al., 2024): on Llama-2-7B/WikiText-2, the exact model and dataset DISP-LLM reports its own headline numbers on, LEP beats DISP-LLM’s published perplexity at every converged operating point from 5% to 29% of total model parameters pruned. This advantage does not fully transfer to downstream zero-shot accuracy, however: LEP is at parity with DISP-LLM through roughly 9% of parameters pruned but falls up to 3.0 points behind at the highest sparsity we test, a divergence we report and discuss rather than omit. We also report a case study applying the method ahead of downstream LoRA fine-tuning on CNN/DailyMail summarization: an early result appeared to show pruning *improving* task performance, but a matched dense-vs-pruned fine-tuning control shows this is attributable to fine-tuning itself, not sparsity — what survives is a genuine compression result, a Llama-2-7B model with 26.6% fewer parameters matching fine-tuned dense quality on ROUGE-L at a modest perplexity cost, which we report alongside the caveats that keep it from being a stronger claim. A second, independent case study on SAMSum dialogue summarization surfaces two further genuine findings. First, task-relevant behavioral specialization, producing a genuine, appropriately short summary instead of continuing the dialogue or rambling to the generation-length cap, emerges from pruning alone, before any fine-tuning and without a single gradient update ever touching the frozen model’s own weight values, purely as a consequence of which neurons a separately-trained selector network chooses to keep. Second, once both arms are fine-tuned under matched budgets, the pruned model reaches 94.7% of dense fine-tuned quality on ROUGE-L, a real, bootstrap-confirmed 2.25-point gap rather than

the CNN/DailyMail case’s statistical tie, while retaining only 48.2% of total parameters (80.6% of feed-forward neurons removed) — an honestly quantified compression-for-quality trade, not a free one.

1 Introduction

Large language models are expensive to serve. Structured sparsity — removing whole neurons, attention heads, or channels rather than scattering unstructured zeros through a weight matrix — is the form of compression that actually converts into wall-clock and memory savings on commodity hardware, unlike unstructured magnitude pruning, which typically needs specialized sparse kernels to realize any speedup at all (Frantar & Alistarh, 2023; Sun et al., 2023).

Most structured-pruning methods for LLMs share two properties. First, the importance signal is a **hand-designed, local rule**: weight magnitude, activation magnitude (Hu et al., 2016), a first-order Taylor estimate (Molchanov et al., 2017; Ma et al., 2023), or a closed-form layerwise reconstruction objective (Frantar & Alistarh, 2023; Sun et al., 2023; Guo et al., 2025). Second, that signal is computed **independently per neuron**, with cross-layer interaction handled only implicitly, if at all, through a uniform or heuristically-tuned per-layer budget. A small number of recent methods instead *learn* the scoring function via gradient descent — most notably DISP-LLM (Gao et al., 2024), which trains a hypernetwork to set per-layer widths — but this remains rare, and the learned methods that do exist do not condition their scoring jointly on the full weight vector of every neuron together with an explicit, depth-aware sharing mechanism between layers.

We take the position that a trained network’s redundancy structure is already legible in its weights, and that the right question is not “which neurons have small activations on this calibration set” but “what does the weight matrix itself say about how compressible this network is.” LEP answers this with a hypernetwork that (a) scores each neuron from its own weight row via a shared, permutation-respecting encoder, and (b) shares that information across the depth axis via a bidirectional LSTM, so a layer’s implied redundancy can be informed by every other layer’s before gates are decided. The pruner is trained end-to-end against the frozen base model’s own loss, with a single sparsity penalty λ as the only knob — no external labels beyond next-token prediction, no calibration set beyond the pruner’s own training batches, and no architecture-specific tuning needed to port the same code across four different base models spanning three orders of magnitude in parameter count.

Contributions.

1. A weight-conditioned hypernetwork architecture (row-encoder + bidirectional cross-layer context) for structured neuron pruning of frozen networks, with a proved permutation-equivariance guarantee (§3.3).
2. A toy-scale pilot connecting the method to the Lottery Ticket Hypothesis (Frankle & Carbin, 2019): weights alone, without data or gradients or retraining, are sufficient to discover a highly compressible subnetwork.
3. Clean, reconciled sparsity-perplexity Pareto frontiers on GPT-2 small and OPT-125M, an efficiency metric characterizing the frontier’s best operating point, and a $\sim 6\times$ perplexity-cost improvement over a matched-sparsity activation-magnitude baseline.
4. A direct, head-to-head comparison against a recent trained-hypernetwork baseline (DISP-LLM) on the exact model and dataset it reports its own numbers on — Llama-2-7B, WikiText-

2 — where LEP beats the published curve at every converged operating point, together with an honest diagnosis of the one operating point where it currently does not.

2 Related Work

Lottery Ticket Hypothesis. Frankle & Carbin (2019) show that dense, randomly-initialized networks contain sparse subnetworks that, trained in isolation from the *same* initialization, match the dense network’s accuracy. Liu et al. (2019a) complicate this by showing the pruned *architecture*, not the specific inherited weights, often carries most of the value. Our toy-scale pilot (§4) connects to this literature but differs mechanically: we never retrain or rewind — the mask is read directly off the trained weights by a learned scorer — and we care about frozen-model *inference-time* sparsity, not *trainability from a lottery-ticket initialization*.

Classical structured-pruning scoring rules. Magnitude pruning (Han et al., 2015) remains the strongest simple baseline; activation-magnitude/APoZ scoring (Hu et al., 2016) and first-order Taylor-expansion importance (Molchanov et al., 2017; Ma et al., 2023) are the next step up in sophistication. All three score neurons independently, with no learned cross-layer interaction — this is the baseline family §5.4’s comparison is measured against.

One-shot / calibration-based LLM pruning. SparseGPT (Frantar & Alistarh, 2023) and Wanda (Sun et al., 2023) prune in a single forward pass via a closed-form layerwise reconstruction objective, with no training loop at all — SparseGPT reports up to 60% *unstructured* sparsity on OPT-175B at “comparable” perplexity (its own reported 50%-sparsity point is, in fact, slightly *below* the dense baseline: 8.35 $\rightarrow$ 8.21 ppl on raw-WikiText2), and Wanda matches this quality at roughly 300 $\times$ lower scoring cost by dropping the weight-update step entirely. LLM-Pruner (Ma et al., 2023) scores weights via a first-order Taylor expansion and recovers performance with LoRA fine-tuning; FLAP (An et al., 2023) is retraining-free, scoring by input-fluctuation stability. SlimLLM (Guo et al., 2025) evaluates channel and attention-head importance holistically (Pearson similarity between original and pruned-layer outputs) rather than by aggregating per-weight scores, and recovers performance with a fast linear-regression repair of the output matrix; its headline “98.7% of original performance retained” figure is *downstream commonsense-reasoning accuracy with post-pruning LoRA tuning applied* — the directly comparable, tuning-free number in the same table is 96.8%. GISP (Wang et al., 2025) globally (not per-layer) prunes attention heads and MLP channels via an iterative, block-normalized first-order importance signal, and is, to our knowledge, the only prior structured-LLM- pruning work reporting numbers on both Llama-3-8B and Mistral-7B — the exact models this paper also validates on (§5). SDS (Li et al., 2024) reconstructs weights between two one-shot semi-structured (2:4) prune passes and reports near-full dense recovery at that sparsity pattern on OPT-125M. None of these methods train a scoring network; all are calibration-based, whether closed-form or iterative.

Learned, trained-mask pruning. CoFi (Xia et al., 2022) and L_0 regularization (Louizos et al., 2017) jointly learn coarse- and fine-grained binary masks via a relaxed L_0 penalty, evaluated on BERT-scale classification tasks, not autoregressive LLMs. Closest to our own method is **DISP-LLM** (Gao et al., 2024), which trains a hypernetwork (a GRU $\rightarrow$ LayerNorm $\rightarrow$ GeLU $\rightarrow$ Linear stack) via gradient descent (using the ReinMax straight-through estimator) to set per-layer widths for both attention and MLP structures jointly, without updating the base model’s weights — architecturally the closest analogue to LEP in the literature we are aware of. Two differences are load-bearing

for the comparison in §6.3: DISP-LLM’s training objective is the pruned model’s own language-modeling loss alone (no paired forward pass through the unpruned model), and it prunes across attention *and* MLP structures jointly, a strictly broader scope than LEP’s FFN-only pruning. DISP-LLM also trains for a fixed 10,000 iterations per (model, target-sparsity) pair regardless of difficulty, versus LEP’s convergence-based, difficulty-adaptive stopping rule (§3.6) — see §6.3 for the resulting compute comparison. Structured channel pruning via a trained meta-network has precedent in the CNN literature (Network Slimming, Liu et al., 2017; MetaPruning, Liu et al., 2019b) but is not evaluated here; our scope is LLM feed-forward blocks.

Reinforcement learning and hard-budget formulations. We considered and rejected two alternative formulations during development. A sequential, RL-based pruning policy is provably dominated by direct differentiable optimization on a frozen model: because the per-step reward telescopes, the episode return depends only on the *final* set of kept neurons, not the order they were chosen, so the induced MDP has path-independent return and no credit-assignment structure for RL to exploit. A hard top- K budget with a straight-through estimator gives gradient only to neurons near the moving threshold and is both less stable and substantially weaker at matched accuracy than the soft λ -penalty formulation used throughout this paper. Neither is discussed further.

3 Method

3.1 Problem setup

Let f_θ be a frozen, pretrained language model; θ is never updated, before or after pruning. A pruner g_ϕ reads f_θ ’s own weight matrices and outputs per-neuron binary gates m . The objective is to find ϕ minimizing task loss on the gated model $f_\theta(x; m)$ subject to a sparsity penalty on m . The deliverable is the pair (θ, m) — a smaller effective model — not a better θ . This paper does not claim, anywhere, that pruning improves the base model; §7.2 explains why an early, now-corrected result once looked like it did.

3.2 Pruner architecture

For each prunable layer ℓ , let $W_\ell \in \mathbb{R}^{n_\ell \times d_\ell}$ be the weight matrix whose rows are treated as one token per neuron. A **row encoder** — a shared 2-layer MLP, $\text{Linear}(d_\ell, \text{embed_dim}) \rightarrow \text{ReLU} \rightarrow \text{Linear}(\text{embed_dim}, 1)$, applied identically to every row — maps each neuron’s weight vector to a scalar logit. The output bias is initialized to +2.0, so every gate starts near-fully-open: pruning is learned in, not defaulted. In parallel, each layer’s weight matrix is summarized by a mean over the neuron axis and projected to a shared hidden size, and the resulting per-layer embeddings are run through a **bidirectional LSTM** over the depth sequence. The LSTM’s output, passed through $\tanh$ (bounding it to $(-1, 1)$ so it can *modulate* but never override the per-neuron logit) and a small linear head, is added back to every neuron’s logit in that layer as a cross-layer context bias. Gates are produced with a straight-through estimator: hard threshold at 0.5 in the forward pass, sigmoid gradient in the backward pass. The same architecture and hyperparameters ($\text{embed_dim} = 64$, $\text{lstm_hidden} = 128$) are reused, unmodified beyond the weight-extraction plumbing, across every base model in this paper — MLP hidden layers, a CIFAR conv-net’s FC head, GPT-2’s Conv1D FFN, OPT-125M’s `nn.Linear` FFN, and the SwiGLU FFN blocks of Mistral-7B and Llama-2-7B (§5).

3.3 Permutation equivariance

Neurons within one layer have no canonical order: any permutation of a layer’s rows, with the matching permutation applied downstream, computes an identical function. A pruning criterion that is not invariant to this symmetry is fitting an artifact of storage order rather than the network’s actual redundancy structure. The row encoder is exactly invariant to this symmetry by construction — it is one shared function $h : \mathbb{R}^{d_\ell} \rightarrow \mathbb{R}$ applied independently per row, so permuting the rows just permutes which score goes with which row, and the layer-context summary (mean over the neuron axis) is likewise permutation-invariant.

Lemma 1 (Permutation equivariance). *For any permutation matrix P acting on layer ℓ ’s neuron axis, $g_\phi(PW_\ell) = P \cdot g_\phi(W_\ell)$.*

Proof sketch. The row encoder h is applied row-wise and identically to every row, so $h(PW_\ell) = Ph(W_\ell)$ directly. The layer-context term depends on W_ℓ only through $\text{mean}_i(W_{\ell,i})$, which is invariant to any permutation of the row index i ; it therefore contributes an identical additive bias to every row regardless of P . The gate function is a fixed, row-wise nonlinearity of the sum of these two terms, so it commutes with P as well. $\square$

This is a structural guarantee, not a trained-in behavior. It also rules out one candidate explanation for a limitation we discuss in §7.3 (a trained pruner does not transfer across independently-trained networks): the failure is not a missing permutation invariance, since the invariance is already exact.

3.4 Why a sequence model over depth, and not over neurons

Depth, unlike within-layer neuron order, is a genuine, non-arbitrary axis: layer ℓ ’s output literally feeds layer $\ell+1$. A sequence model over the depth axis is therefore the structurally appropriate choice, in direct contrast to a sequence model over neurons-within-a-layer, which would impose a false order onto a permutation-symmetric set — exactly the mistake §3.3 rules out for the row encoder. The bidirectional LSTM lets a layer’s implied sparsity budget depend on the redundancy profile of every other layer in both directions: a layer with unusually concentrated redundancy can inform neighboring layers that they can be pruned harder without creating a capacity bottleneck. Our toy-scale pilot (§4) motivates this design directly — prunability tracks layer width, and a pruner that cannot see across layers cannot discover that asymmetry.

3.5 Training objective

$$\mathcal{L}(\phi) = (\text{CE}_{\text{pruned}} - \text{CE}_{\text{orig}}) + \lambda \cdot \frac{1}{L} \sum_{\ell=1}^L \bar{g}_\ell, \quad (1)$$

where CE_{orig} and $\text{CE}_{\text{pruned}}$ are the frozen base model’s cross-entropy loss on the same batch with all gates open and with the current gates applied, respectively, and $\bar{g}_\ell$ is the mean gate value in layer ℓ (so the penalty decreases as sparsity increases). Only ϕ receives gradients; the frozen model’s forward pass runs twice per step, once ungated and once gated. λ is the sole knob controlling the sparsity/accuracy operating point and is swept externally (§3.7).

We flag one open methodological question rather than resolving it silently: as written, the loss has no floor at zero, so gradient descent is free to push $\text{CE}_{\text{pruned}}$ *below* CE_{orig} if the training distribution allows it. §7.2 discusses exactly this happening early in this project’s OPT-125M experiments, why it does not reflect a genuine capability improvement, and why we report it as a

cautionary methodological finding rather than retroactively re-running every checkpoint against a floored objective, $\max(\text{CE}_{\text{pruned}} - \text{CE}_{\text{orig}}, 0)$. All results in this paper use the unfloored objective as written above.

3.6 Convergence-based stopping and plateau-triggered learning-rate decay

Rather than a fixed step budget — which we found to be badly and non-uniformly mismatched to actual convergence across λ and models — training stops via a block-mean flatness check: every **check_every** steps, the mean gate value per layer over the trailing block is compared against the same statistic at **window** prior checkpoints, and training is declared converged if every layer’s block-mean is within a relative/absolute tolerance across the whole window. A raw convergence signal is not trusted immediately: it triggers a cosine learning-rate decay ($\text{lr}_0 \rightarrow \text{lr}_{\min}$ over a fixed window), and only a *second* flatness check, evaluated once the decay window has elapsed, confirms convergence. If the reconfirmation fails, the run holds at $\text{lr}_{\min}$ and continues — this correctly distinguishes a genuine plateau from a noisy trajectory that happened to look flat over one window (§6.4 discusses one operating point where this mechanism still does not resolve a genuine, persistent instability).

3.7 λ sweep, Pareto curve, and efficiency metric

λ is swept on a grid (typically log-spaced) per base model, 2 seeds per point at GPT-2/OPT-125M scale where compute allowed, 1 seed at 7B scale. We report the full **Pareto curve** — percent of neurons (or, where directly comparable to prior work, percent of total model parameters) pruned versus pruned-model perplexity, against the unpruned baseline as a reference line — alongside a single-number **efficiency metric**,

$$\text{efficiency}(\lambda) = \frac{\% \text{ pruned}}{\exp(\Delta \text{CE})}, \quad \Delta \text{CE} = \ln \left(\frac{\text{ppl}_{\text{pruned}}}{\text{ppl}_{\text{orig}}} \right), \quad (2)$$

which converts the nats-scale cost back to a perplexity-ratio scale so it is comparable across models and datasets with different baseline cross-entropy. We use this to identify a recommended operating point, while reporting the full curve alongside it since the efficiency peak is typically a broad plateau, not a sharp optimum.

4 Toy-Scale Pilot: Motivating Evidence

Before scaling to language models, we ran a pilot on small MLPs (MNIST) and a convolutional network’s fully-connected head (CIFAR-10). We report this section qualitatively only — no sparsity percentages, accuracy-drop figures, or weight-count numbers — since these toy-scale results exist to motivate the hypotheses tested for real at LLM scale (§5-§6), not as this paper’s empirical content.

Four qualitative regularities motivated the design in §3. First, **weights alone are sufficient to find a good mask**: no data, no gradients, and no retraining cycle are needed to discover a highly compressible subnetwork directly from a trained model’s weight matrices — a stronger claim than the classical Lottery Ticket Hypothesis’s iterative reinitialize-and-retrain procedure (Frankle & Carbin, 2019), since here the mask is read off directly. Second, **the discovered mask is a near-deterministic property of the trained weights**, not an artifact of the pruner’s own random initialization or training noise — supporting the reading that “prunability” is a real, measurable property of the base network rather than measurement noise. Third, **the learned pruner dominates classical scoring rules at every accuracy budget tested**, because it conditions on

the full weight vector rather than a hand-designed scalar summary — the toy-scale precedent for §5.4’s LLM-scale baseline comparison. Fourth, **prunability is governed by layer width, not depth or raw parameter count**: wide layers concentrate redundancy, while depth behaves as load-bearing capacity that resists pruning — directly motivating §3.4’s cross-layer context design, since a pruner that cannot see across layers cannot discover this asymmetry.

5 LLM-Scale Experiments

5.1 Setup

We prune the FFN intermediate neurons of frozen GPT-2 small and OPT-125M (both 12 layers, hidden size 768, architecturally compute-equivalent), training and evaluating on WikiText-2 with a standard sliding-window cross-entropy evaluation protocol (2048-token windows, 1024-token stride), using the same pruner architecture and hyperparameters as the toy-scale pilot with no per-model tuning beyond the weight-extraction plumbing.

5.2 GPT-2 small




Figure 1: Sparsity-perplexity Pareto frontiers for GPT-2 small and OPT-125M, WikiText-2, frozen-model FFN pruning.

A $\lambda = 0.01 \rightarrow 3.2$ sweep (2 seeds per point) gives a clean, monotonically increasing cost curve across the entire range — no free-lunch region, no crossover with the dense baseline (dense ppl = 25.109). Table 1 reports the mean over seeds at each λ ; peak efficiency (§3.7) falls at $\lambda = 1.35$ (49.16% pruned, ppl 44.075 ± 1.94), with a broad plateau from roughly $\lambda = 0.75$ to $\lambda = 2.4$.

5.3 OPT-125M

A 9-point, convergence-based sweep (dense ppl = 23.941; single seed, convergence-based stopping per §3.6 rather than a fixed step budget) gives a qualitatively similar shape, with one wrinkle worth flagging plainly rather than hiding: at the three lowest λ values (≤ 0.1), the measured pruned-model perplexity is at or slightly *below* the dense baseline. We do not read this as pruning improving the model — §7.2 explains the mechanism (an unfloored loss term combined with WikiText-2’s own correlated train/test split can produce exactly this signature, and it does not survive an out-of-

Table 1: GPT-2 small, WikiText-2, mean over 2 seeds (dense ppl = 25.109).

λ	% pruned	pruned ppl	Δ ppl
0.55	25.01%	30.253	+5.14
0.75	32.00%	35.766	+10.66
1.0	39.05%	40.064	+14.96
1.35	49.16%	44.075	+18.97 (peak efficiency)
1.8	56.71%	55.862	+30.75
2.4	67.99%	63.013	+37.90
3.2	76.50%	73.201	+48.09

domain check) and why we still report the raw numbers rather than adjust them. The genuine, monotonic cost region begins by $\lambda = 0.2$.

Table 2: OPT-125M, WikiText-2, single seed, convergence-based stopping (dense ppl = 23.941).

λ	% pruned	pruned ppl	Δ ppl
0.01	7.41%	23.603	-0.34 (see §7.2)
0.05	14.75%	22.723	-1.22 (see §7.2)
0.1	19.56%	22.648	-1.29 (see §7.2)
0.2	22.81%	24.422	+0.48
0.25	25.66%	24.027	+0.09
0.3	27.79%	25.160	+1.22
0.4	31.42%	26.050	+2.11
0.8	41.62%	28.632	+4.69
1.6	63.39%	38.737	+14.80

5.4 Baseline comparison

At OPT-125M’s $\lambda = 0.75$ operating point (42.56% pruned; this specific point comes from an earlier fixed-step sweep, not the convergence-based sweep in Table 2 — flagged for reproducibility, not a discrepancy in the result), we compare against a standard activation-magnitude baseline: mean post-ReLU activation over 50 calibration batches of WikiText-2 train, global magnitude threshold picked to match the trained pruner’s exact neuron count. Both are evaluated under the identical protocol.

Table 3: LEP vs. activation-magnitude baseline, OPT-125M, 42.56% pruned (matched sparsity).

method	orig ppl	pruned ppl	Δ ppl
LEP (this work)	23.941	27.086	+3.145
Activation-magnitude baseline	23.944	43.818	+19.873

LEP’s perplexity cost is $19.873/3.145 \approx 6.3\times$ smaller than the activation-magnitude baseline at exactly matched sparsity. Qualitatively, the baseline’s failure mode concentrates almost all of its cuts in a few middle layers (down to $\sim 13\%$ kept in its worst layer), while LEP’s per-layer allocation

stays comparatively flat across depth — consistent with §3.4’s cross-layer-context design mattering in practice, not just in theory.

6 Scaling to 7B: Validation Against Published Baselines

6.1 Setup

We port the identical pruner architecture and training regime to Mistral-7B (Apache 2.0, un-gated) and Llama-2-7B (gate-licensed; access requested separately). Both use a SwiGLU feed-forward block — $\text{down_proj}(\text{SiLU}(\text{gate_proj}(x)) \odot \text{up_proj}(x))$ — architecturally different from GPT-2/OPT-125M’s single-matrix FFN. The gate hook attaches to `down_proj`’s input (the direct SwiGLU analogue of GPT-2’s `c_proj`/OPT’s `fc2`), and each neuron’s row-encoder input is the concatenation of its `gate_proj` and `up_proj` rows. Base-model weights are stored in bf16; only the small extracted weight slice fed to the pruner’s own encoder is upcast to fp32. No other architectural change was needed.

6.2 Mistral-7B

On C4 (dense ppl = 8.325), a $\lambda = 0.01 \rightarrow 0.4$ sweep (one point, $\lambda = 0.8$, did not converge and is excluded) shows the same free-region-then-monotonic-cost shape seen at every smaller scale: near-zero or slightly negative cost at $\lambda \leq 0.05$ (2.84% pruned, +0.02 ppl), rising to +0.72 ppl at 18.49% pruned. On WikiText-2 (dense ppl = 4.740), the same model shows a much flatter absolute cost curve (e.g. +0.58 ppl at 18.92% pruned) — consistent with Mistral-7B already being well-calibrated to clean Wikipedia text, leaving less genuine redundancy to find there specifically, exactly the concern we flagged before running this sweep.

6.3 Llama-2-7B vs. DISP-LLM

This is the paper’s most direct external validation: DISP-LLM (Gao et al., 2024) reports its own headline numbers on this exact model and dataset — Llama-2-7B, WikiText-2, no base-weight update — so this comparison has no cross-model confound of the kind that would complicate, say, a Mistral-7B vs. Llama-2-7B comparison.

We convert LEP’s FFN-neuron-pruned percentages to total-model-parameter percentages (Llama-2-7B’s FFN blocks are 64.24% of total parameters) to match DISP-LLM’s reporting convention, and compare against DISP-LLM’s own Table 1, linearly interpolated to our exact operating points.

LEP beats DISP-LLM’s interpolated curve at all 8 converged points, by a wide margin at low sparsity, narrowing to an effective tie by 28.8% total-param pruned. The one point that loses (32.58%) never converged (§6.4) and is not treated as a real comparison point.

On compute: DISP-LLM trains a fixed 10,000 iterations per (model, target-sparsity) pair regardless of difficulty — 2.41 hours on 2×A100 per point for the LLaMA/LLaMA-2 7B family (Gao et al., 2024, Table 6), i.e. ≈ 4.82 A100-GPU-hours per operating point, with no separate unpruned-model forward pass per step. LEP’s convergence-based stopping is difficulty-adaptive: our 8 converged Llama-2-7B points ranged from 553 seconds ($\lambda = 0.05$) to 5,211 seconds ($\lambda = 1.0$) on a single H100, with the harder, higher-sparsity points naturally costing more rather than every point paying a fixed budget regardless of how quickly it actually settles.




Figure 2: LEP vs. DISP-LLM, Llama-2-7B, WikiText-2, no base-weight update. LEP beats DISP-LLM’s published curve at every converged operating point; the one point that does not ($\lambda = 1.4$) never converged and is excluded from the headline claim (§6.4).

Table 4: LEP vs. DISP-LLM, Llama-2-7B, WikiText-2, no weight update (dense ppl: ours = 4.903, DISP-LLM = 5.12).

% total params pruned	LEP ppl	DISP-LLM ppl (interp.)	LEP – DISP-LLM	converged?
5.39%	4.933	5.384	−0.451	yes
6.57%	4.992	5.442	−0.450	yes
8.89%	5.118	5.556	−0.438	yes
11.58%	5.275	5.687	−0.412	yes
15.05%	5.591	5.857	−0.266	yes
19.79%	5.865	6.090	−0.225	yes
24.38%	6.409	6.428	−0.019	yes
28.77%	6.774	6.758	+0.016	yes
32.58%	7.669	7.175	+0.494	no

6.4 An honestly diagnosed failure mode

The $\lambda = 1.4$ operating point (32.58% total-param pruned) never converges within an 18,000-step safety cap — not a near-miss, a genuine, persistent instability. In the final 20 diagnostic checkpoints (the last 3,800 steps of the run), the fraction pruned still swings across an 8.2 percentage-point band with almost the same noise level as the *first* 20 checkpoints, and individual layers still move by up to 47.8 percentage points between consecutive 200-step checkpoints at the very end of training — an oscillating equilibrium, not a slow monotonic approach to a plateau that more steps would resolve.

We considered pruner capacity as the likely fix and tested it directly against this project’s own prior evidence rather than assuming it: an 8-point matched ablation at $0.5\times$, $1\times$, and $2.26\times$ pruner capacity (GPT-2 and OPT-125M, a separate corpus) found *no* monotonic effect of capacity on percent pruned, perplexity, or convergence speed. Capacity controls the row-encoder/BiLSTM’s representational expressiveness; it has no obvious mechanistic connection to the straight-through gate’s threshold-oscillation dynamics under a fixed learning rate. Our current best explanation is a learning-rate/ λ mismatch: at this operating point the sparsity penalty term ($\lambda \cdot \bar{g} \approx 0.66\text{--}0.8$) is comparable in magnitude to the cross-entropy cost term ($\approx 0.37\text{--}0.46$), while every λ in every sweep in this paper uses the same fixed initial learning rate — a plausible recipe for sustained oscillation around a ridge in the loss landscape rather than convergence to a fixed point, since the plateau-triggered decay mechanism (§3.6) never gets a chance to intervene if the raw convergence check never fires even once. We report this as an open, unresolved methods gap rather than retrying with a larger step budget and hoping.

6.5 Downstream zero-shot accuracy vs. DISP-LLM: the perplexity advantage does not transfer

§6.3 compares perplexity; comparable papers (DISP-LLM, SlimLLM, LLM-Pruner, FLAP, GISP) also report zero-shot downstream accuracy, so we ran the same 8 converged Llama-2-7B/WikiText-2 checkpoints plus dense through lm-evaluation-harness (PIQA, HellaSwag, WinoGrande, ARC-easy, ARC-challenge, BoolQ, OpenBookQA, 0-shot). Restricting to the 5-task subset DISP-LLM’s own Table 3 reports (WinoGrande acc; HellaSwag/ARC-e/ARC-c/PIQA acc-norm — matching metric convention exactly) and comparing against DISP-LLM’s own published numbers:

Table 5: LEP vs. DISP-LLM zero-shot downstream accuracy, Llama-2-7B, WikiText-2, no weight update (5-task average).

% total params pruned	LEP avg acc	DISP-LLM avg acc (interp.)	LEP – DISP-LLM
0% (dense)	68.46	68.99	−0.53
5.39%	66.59	67.03	−0.44
6.57%	66.91	66.61	+0.30
8.89%	66.28	65.76	+0.52
11.58%	64.37	64.79	−0.42
15.05%	62.87	63.53	−0.66
19.79%	60.29	61.81	−1.52
24.38%	59.28	60.14	−0.86
28.77%	55.51	58.55	−3.04

(% total params pruned uses each checkpoint’s own training-time “final % FFN neurons pruned”

— the same figures underlying Table 4 — not a fresh gate reconstruction at eval time; see the note below.)

Unlike Table 4’s ppl comparison, DISP-LLM’s downstream table reports only *two* real Llama-2-7B operating points (30%, 50%, vs. five for ppl), so the “interpolated” column here is a much cruder straight line between their dense point and their single 30% point. Our 28.77% point sits almost exactly at their measured 30% checkpoint, so a direct, interpolation-free comparison is available and more defensible: against DISP-LLM’s own actually-measured 30% value (58.10) — at a slightly *more* aggressive pruning ratio than our 28.77%, i.e. charitable to DISP-LLM, not to us — LEP is still behind by **2.59 points**.

Reconstruction note. A fresh eval-time gate reconstruction from each checkpoint’s saved pruner weights does not reproduce the training-time “final % FFN neurons pruned” figure exactly — negligibly so for 7 of 8 checkpoints ($< 0.05\text{pp}$), but 44.23% reconstructed vs. 44.78% training-logged at $\lambda = 1.0$ (0.55pp, the largest gap observed), despite the reconstructed checkpoint’s perplexity matching the training log exactly. The accuracy numbers above come from whatever mask the reconstruction actually produced, so a small mismatch between “the mask we’re reporting accuracy for” and “the mask’s training-time self-reported sparsity” exists at $\lambda = 1.0$ specifically; the table’s x-axis label instead uses the training-time figure for consistency with Table 4. Mechanism unconfirmed (candidates: an eval-mode vs. training-mode difference somewhere in the pruner forward pass, or borderline gate values near the STE threshold flipping under a different numerical path) — flagged as an open reproducibility question, not resolved here.

The shape of this result is the opposite of Table 4’s: there, LEP starts ahead of DISP-LLM and the margin narrows toward high sparsity; here, LEP starts at parity (within $\pm 0.5\text{pp}$, consistent with noise on a 7-task harness) through roughly 9% total-param pruned, and the deficit *widens* as pruning increases, reaching -3.04pp (interpolated) / -2.59pp (direct) at our most aggressive converged point. This is a genuine, unresolved divergence between the two metrics, not a result to smooth over — a perplexity advantage over a published baseline does not, by itself, guarantee an equivalent downstream-accuracy advantage. Our current best (untested) explanation: perplexity is a smooth, dense, next-token-averaged metric measured on text distributionally close to WikiText-2 — the same corpus the pruner is trained against — while the downstream suite is discrete, out-of-domain multiple-choice/QA, and may depend disproportionately on specific circuits an in-domain-ppl-optimized mask does not protect. We have not tested this directly and report it as an open question (§7.5), not a conclusion.

6.6 Case study: pruning plus LoRA fine-tuning on CNN/DailyMail summarization

Every result so far in this paper leaves the base model’s weights untouched, before and after pruning (§3.1). This section is a deliberate, clearly-scoped exception: we ask whether a pruned model recovers task performance under *subsequent* LoRA fine-tuning, and whether pruning contributes anything a plain fine-tune of the dense model does not already give for free. We report it as a self-contained case study, not part of the paper’s main compression claim.

Motivation. An early signal (single run, $\lambda = 0.1$, non-converged, Llama-2-7B, CNN/DailyMail, target-only cross-entropy) appeared to show a pruned model *beating* the frozen dense model on both held-out perplexity ($3.750 \rightarrow 2.881$) and generation-quality ROUGE (rouge-L $13.31 \rightarrow 18.71$). This is the same qualitative shape as an earlier, since-largely-retracted result on OPT-125M/WikiText-2 (§7.2) — but here the training path was clean of that result’s tokenization confound, and the

metric was a real generation task, not in-domain perplexity alone. It was flagged as suggestive, not conclusive: the dense baseline being compared against was zero-shot (no fine-tuning at all), so an apparent “pruning helps” result could just as easily be “any adaptation helps, and this run happened to include one.”

Matched-control design. We resolve this directly: fine-tune both a pruned model ($\lambda = 0.3$, 41.36% of FFN neurons removed, 4,948,201,472 of 6,738,415,616 params retained — a 26.6% total-parameter reduction) and the dense model with LoRA (rank 16, $\alpha = 32$, dropout 0.05, all seven attention/MLP projections, cosine schedule, 3% warmup, weight decay 0), under identical budgets: 1,122 optimizer steps, global effective batch 128, 143,616 examples seen (50% of the 287,113-example CNN/DailyMail train split), identical seed. ROUGE is scored by greedy generation over the same 3,000-example prefix of the test split (26% of the 11,490-example test set) for every arm. Every hyperparameter and scale setting is asserted equal between arms programmatically before any comparison is reported.

Table 6: Pruned vs. dense LoRA fine-tuning, Llama-2-7B, CNN/DailyMail (test, $N = 3,000$; ppl on full 11,490-example test split).

arm	params	ppl	rouge-1	rouge-2	rouge-L
dense, zero-shot	6,738,415,616	3.750	20.53	8.74	13.78
dense + LoRA	6,738,415,616	2.774	34.41	14.90	22.99
pruned, surgery only	4,948,201,472	3.126	33.45	14.16	22.27
pruned + LoRA	4,948,201,472	2.979	34.31	14.93	22.93

Result. Pruned+LoRA reaches rouge-L 22.93 against dense+LoRA’s 22.99 — a 0.06-point gap, well inside plausible noise. Perplexity does not tell the same story: dense+LoRA reaches 2.774 against pruned+LoRA’s 2.979, a real 0.205-point gap that holds at every intermediate checkpoint we logged (steps 250/500/750/1000/1122), with dense already ahead of pruned’s final value by its own step-500 checkpoint.

This directly answers the motivating question, and against the pruning-helps reading: **the apparent pruned-beats-dense pattern was fine-tuning, not sparsity.** The dense arm moves rouge-L from 13.78 to 22.99 (+9.21pp) on LoRA alone, with no pruning anywhere in that arm; the pruned arm reaches essentially the same place. What survives is a compression claim, not a specialization claim: a model with 26.6% fewer parameters matches fine-tuned dense quality on rouge-L, at a real but modest perplexity cost. We did not attempt to characterize where this trade-off holds across other sparsity levels, only that it holds at this one.

Caveats. We report several limitations of this comparison explicitly rather than let the table’s precision overstate its rigor. (1) *No confidence interval*: per-example score pairs were not retained, so the 0.06pp rouge-L gap is a bare point estimate over 3,000 examples, not a statistically separated result — we cannot rule out that the true gap is larger in either direction. (2) *Single seed, single sparsity level* ($\lambda = 0.3$): we do not know whether iso-quality parity holds at other pruning ratios. (3) *Non-converged pruner*: the pruning checkpoint underlying both pruned arms never fully converged in training (§3.6’s block-mean flatness check did not fire before a fixed step cap), an inherited caveat from the motivating run. (4) *Weak zero-shot baseline, unverified*: the dense zero-shot rouge-L (13.78) is low for a non-instruction-tuned model completing a bare prompt; we did

not independently verify this isn’t itself an artifact of prompt formatting, though this affects our reading of the *motivating* result only, not the matched-control table, since both LoRA arms start from and are scored under identical conditions. (5) *Sample-size correction*: the earlier motivating run’s rouge-L numbers used a 300-example evaluation prefix rather than this section’s 3,000; we subsequently found the smaller sample understates pruned-model rouge-L specifically by roughly 3 percentage points while leaving dense-model scores essentially unchanged, meaning the motivating run’s headline gap was, if anything, under- rather than over-stated — the matched-control table above is unaffected, since it always used $N = 3,000$ for every arm.

This section’s scope is deliberately narrow: it shows pruning-then-fine-tune is a viable path to a smaller, task-adapted model at (here) no measurable rouge-L cost, not that pruning improves anything, and not that this trade-off generalizes beyond the one operating point tested.

6.7 Case study: pruning plus LoRA fine-tuning on SAMSum dialogue summarization

We repeat §6.6’s design on a second, independent summarization dataset, SAMSum (dialogue $\rightarrow$ short summary), at the most aggressive operating point in an 8- λ Llama-2-7B/SAMSum sweep ($\lambda = 1.6$). This case study surfaces two results we report as genuine findings of the paper, not as a repeat of §6.6’s conclusion.

Motivation and a methodological fix. The raw sweep showed pruning beating the zero-shot dense model on rouge-L at every λ , including 80%+ sparsity — implausible as a quality claim, and diagnosed as a generation-length artifact: SAMSum references are short (~ 28 tokens), but nothing in the training pipeline ever taught the model to stop, so dense generation rambled to the 96-token cap while continuing the dialogue or emitting meta-commentary rather than summarizing, corrupting rouge-L’s precision term. We fixed this at the source rather than only at evaluation time: we appended an end-of-sequence token to the supervised target span both in the pruner’s own CE-delta training objective (§3.5) and in the downstream LoRA fine-tune, applied identically to every arm below. An earlier attempt that fixed only the LoRA stage, leaving the pruner itself trained without this signal, showed the same qualitative pattern but confounded by residual rambling in the pruned arm’s own pre-fine-tuning generations; we disregard that run and report only the fully-consistent pipeline here.

Matched-control design. We fine-tune both a pruned model ($\lambda = 1.6$, 80.57% of FFN neurons removed, 3,250,896,896 of 6,738,415,616 params retained — a 51.76% total-parameter reduction) and the dense model with LoRA, identical configuration to §6.6 (rank 16, $\alpha = 32$, dropout 0.05, all attention/MLP projections, cosine schedule, 3% warmup, weight decay 0), under identical budgets: 345 optimizer steps (3 epochs over SAMSum’s 14,731-example train split — SAMSum is roughly $20\times$ smaller than CNN/DailyMail, so neither §6.6’s step count nor its “50% of the split” rule transfers), global effective batch 128, identical seed. ROUGE is scored by greedy generation over the *entire* 819-example test split for every arm (no prefix-sampling question to manage, unlike §6.6’s caveat (5)), and per-example scores are retained for both arms, enabling a paired bootstrap confidence interval that §6.6’s comparison lacked.

Finding 1: task-relevant behavior emerges from pruning alone, with no gradient ever touching the frozen model’s weights. Before any fine-tuning, the physically-pruned model already summarizes: rouge-L 30.53 against the dense zero-shot model’s 18.77, a +11.76pp gap. Both models are entirely untrained on this task at this point — the only difference is which 80.57%

Table 7: Pruned vs. dense LoRA fine-tuning, Llama-2-7B, SAMSum (full 819-example test split, both ppl and rouge).

arm	params	ppl	rouge-1	rouge-2	rouge-L	mean gen. tok
dense, zero-shot	6,738,415,616	3.978	24.05	7.47	18.77	95.5
dense + LoRA	6,738,415,616	3.802	50.80	26.66	42.22	26.7
pruned, surgery only	3,250,896,896	3.523	37.99	18.60	30.53	18.7
pruned + LoRA	3,250,896,896	4.161	48.77	24.49	39.97	25.9

of FFN neurons the (separately, gradient-trained) pruning network chose to remove. The gap is visible directly in generation length: mean output collapses from 95.5 tokens (dense, rambling to the cap) to 18.7 tokens (pruned, close to the ~ 28 -token reference length) with no mechanism in either model that was ever told to stop early. The same qualitative pattern holds independently in §6.6: surgery-only rouge-L 22.27 against a 13.78 dense zero-shot baseline (Table 6), a +8.49pp gap, on a different task and a different pruner checkpoint. We are precise about the causal claim: the pruning network’s parameters are trained with gradients against labeled task data (§3.5); what does not happen, in either case study, is any gradient touching the frozen base model’s own weight values. Structural subtraction of neurons, decided by a small trained selector, is on its own sufficient to induce a large, task-relevant behavioral change in a frozen 7-billion-parameter model.

Finding 2: compression at a real, quantified quality cost — not free, unlike §6.6.

Fine-tuning both arms under matched budgets narrows the gap substantially but does not close it: pruned+LoRA reaches rouge-L 39.97 against dense+LoRA’s 42.22, a 2.25-point gap. Unlike §6.6, per-example score pairs were retained for both arms here, giving a paired bootstrap confidence interval over the full 819-example test split (10,000 resamples): mean difference -2.30pp , 95% CI $[-3.39, -1.23]\text{pp}$, $p < 0.0001$ — the gap is statistically separated at this sample size, not sampling noise. We report this plainly rather than as parity: the pruned model reaches 94.7% of dense+LoRA’s rouge-L (39.97/42.22) while carrying 48.24% of the total parameters (3,250,896,896 of 6,738,415,616, from 80.57% of FFN neurons removed within the pruned layers — the two percentages are not the same quantity, and only the total-parameter figure should be read as “model size”). This is a genuine, honestly quantified compression-for-quality trade, distinct in kind from §6.6’s tied result precisely because the cost here is real, measured, and does not vanish under a confidence interval.

Caveats. (1) *Single seed, single sparsity level* ($\lambda = 1.6$, the most aggressive point in the sweep): we do not know whether either finding holds at lighter pruning ratios. (2) *Non-converged pruner*: as in §6.6, the checkpoint never triggered §3.6’s block-mean flatness check before the fixed step cap. (3) *The fix itself was load-bearing*: Finding 1 depended on retraining the pruner with an end-of-sequence-aware objective, not just patching evaluation; we disregard our first, pruner-side-unfixed attempt at this comparison rather than report it, on the grounds that its residual generation-length confound made it uninterpretable, not merely imprecise. (4) *Greedy decoding only* (`num_beams=1`): we did not test whether beam search changes either arm’s behavior. (5) As in §6.6, the dense zero-shot baseline’s weak prompt-following is not independently verified to be free of prompt-formatting artifacts, though this affects only the motivating diagnosis, not the matched-control table, since every arm above is scored under identical conditions.

7 Discussion and Limitations

7.1 Scope

This paper reports a compression method for frozen models: a controlled accuracy-for-sparsity trade, characterized via Pareto curves and an efficiency metric. It is not a fine-tuning technique, and at no point do we claim pruning improves the base model’s capability. We state this plainly here because §5.3 and §7.2 report numbers that could otherwise be misread as such a claim. §6.6 is a deliberate, narrowly-scoped exception where the base model *is* fine-tuned after pruning; we keep its claim (iso-quality compression, not improved capability) explicitly separate from the rest of the paper.

7.2 A cautionary tale, not a result

Early in this project, OPT-125M experiments appeared to show pruning *improving* in-domain WikiText-2 perplexity. This turned out to be substantially a tokenization artifact: WikiText-2’s raw format splits one continuous article into thousands of short, blank-heavy “lines,” and a per-line tokenization call with a tokenizer that prepends a beginning-of- sequence token by default (as OPT’s does) scatters thousands of spurious context-resets through what should be one continuous evaluation stream, inflating the *measured baseline* far more than it inflates the pruned model’s score. Once fixed (join each split into one string, tokenize once), the effect shrank substantially but a residual remained at the very lowest λ values, visible in Table 2. Even setting the tokenization bug aside, an unfloored ($\text{CE}_{\text{pruned}} - \text{CE}_{\text{orig}}$) loss term (§3.5) is structurally *capable* of producing this kind of result on any data distribution where train and test are correlated, as WikiText-2’s own split is — and it does not generalize: an out-of-domain check on C4 (§6.2) shows real, monotonic degradation at every λ , with no analogous free region below the noise floor. We report this mechanism explicitly because it is exactly why §7.1’s scope constraint is the right one, not a limitation to apologize for.

7.3 Transfer does not work

A pruner trained on one network does not transfer to a different, even architecturally identical, independently-trained network — each deployment needs its own training run (on the order of 10-90 minutes at the scales in this paper, growing with model size and λ). Since §3.3 establishes that the architecture is already *exactly* permutation-invariant within a layer, the failure is not a missing invariance — it is that raw weight *values* are not directly comparable across independently-trained networks’ weight-space geometry. Activation- or gradient-conditioned inputs, rather than raw weights, are the natural candidate fix; we have not tested this.

7.4 Compute cost honesty

LEP is a trained procedure, not a one-shot calibration rule: roughly 10-25 minutes of GPU time per (λ, seed) operating point at GPT-2/OPT-125M scale, growing to under 10 minutes for easy operating points and up to ~ 90 minutes for hard ones at 7B scale (§6.3), versus one-shot methods like SparseGPT/Wanda, which require minutes total and no training loop. We state this plainly rather than let the Pareto curves imply compression is free; §6.3’s compute comparison against DISP-LLM shows this tradeoff is favorable relative to at least one other trained-hypernetwork method, but not relative to calibration-only baselines.

7.5 Open items

The $\lambda = 1.4$ non-convergence (§6.4) is a genuine, unresolved methods gap — the fixed, λ -independent learning rate used identically across every sweep in this paper is itself a load-bearing, never-revisited default. Every 7B-scale result in this paper is single-seed, unlike the 2-seed protocol used at GPT-2/OPT-125M scale, purely for compute reasons; we do not have a seed-variance estimate at 7B scale and do not claim one. The loss-floor question raised in §3.5 remains open. Pruner-capacity scaling with base-model size (does the minimum viable pruner size grow with the base model, and how) is untested — every experiment in this paper, at every scale, reused the identical pruner configuration. Most consequentially (§6.5): the perplexity advantage over DISP-LLM does not transfer to downstream zero-shot accuracy at high sparsity, and the mechanism behind that divergence is not established — this is the single caveat most likely to change a reader’s overall assessment of the method and should not be read as a footnote.

8 Conclusion

A trained language model’s redundancy structure is legible in its weights. A small, weight-conditioned hypernetwork — a permutation-equivariant per-neuron encoder plus a depth-aware bidirectional context — reads it out directly, without touching the base model’s weights, and produces a clean, controllable sparsity/perplexity Pareto frontier that beats a standard activation-magnitude baseline by roughly $6\times$ at matched sparsity. The same architecture and hyperparameters, unmodified beyond weight-extraction plumbing, port across GPT-2 small, OPT-125M, Mistral-7B, and Llama-2-7B. On the one head-to-head comparison available against a recent trained- hypernetwork baseline in the published literature — DISP-LLM, on the exact model and dataset it reports its own numbers on — this method wins at every converged operating point. We report this result alongside its rough edges: one high-sparsity operating point that does not yet converge, a mechanism we believe explains why, and the open questions that follow from it. The value proposition throughout is a better, more rigorously characterized compression method — not a claim of improved capability, and not a result with its failure modes edited out.

References

- An, Y., Zhao, X., Yu, T., Tang, M., & Wang, J. (2023). FLAP: Fluctuation-based Adaptive Structured Pruning for Large Language Models.
- Frankle, J., & Carbin, M. (2019). The Lottery Ticket Hypothesis: Finding Sparse, Trainable Neural Networks. *ICLR*.
- Frantar, E., & Alistarh, D. (2023). SparseGPT: Massive Language Models Can Be Accurately Pruned in One-Shot. *ICML*.
- Gao, S., Lin, C.-H., Hua, T., Zheng, T., Shen, Y., Jin, H., & Hsu, Y.-C. (2024). DISP-LLM: Dimension-Independent Structural Pruning for Large Language Models. *NeurIPS*.
- Guo, J., Chen, X., Tang, Y., & Wang, Y. (2025). SlimLLM: Accurate Structured Pruning for Large Language Models. *ICML*.
- Han, S., Pool, J., Tran, J., & Dally, W. (2015). Learning both Weights and Connections for Efficient Neural Networks. *NeurIPS*.

- Hu, H., Peng, R., Tai, Y.-W., & Tang, C.-K. (2016). Network Trimming: A Data-Driven Neuron Pruning Approach towards Efficient Deep Architectures.
- Li, G., Zhao, X., Liu, L., Li, Z., Li, D., Tian, L., He, J., Sirasao, A., & Barsoum, E. (2024). Enhancing One-shot Pruned Pre-trained Language Models through Sparse-Dense-Sparse Mechanism (SDS). *arXiv:2408.10473*.
- Liu, Z., Sun, M., Zhou, T., Huang, G., & Darrell, T. (2019). Rethinking the Value of Network Pruning. *ICLR*.
- Liu, Z., Li, J., Shen, Z., Huang, G., Yan, S., & Zhang, C. (2017). Learning Efficient Convolutional Networks through Network Slimming. *ICCV*.
- Liu, Z., Mu, H., Zhang, X., Guo, Z., Yang, X., Cheng, K.-T., & Sun, J. (2019). MetaPruning: Meta Learning for Automatic Neural Network Channel Pruning. *ICCV*.
- Louizos, C., Welling, M., & Kingma, D. P. (2017). Learning Sparse Neural Networks through L_0 Regularization.
- Ma, X., Fang, G., & Wang, X. (2023). LLM-Pruner: On the Structural Pruning of Large Language Models. *NeurIPS*.
- Molchanov, P., Tyree, S., Karras, T., Aila, T., & Kautz, J. (2017). Pruning Convolutional Neural Networks for Resource Efficient Inference. *ICLR*.
- Sun, M., Liu, Z., Bair, A., & Kolter, J. Z. (2023). A Simple and Effective Pruning Approach for Large Language Models (Wanda).
- Wang, Z., Diao, E., Le, Q., Wang, P., Lee, M., Yeh, S., Stupachenko, E. V., Feng, H., & Yang, L. (2025). From Local to Global: Revisiting Structured Pruning Paradigms for Large Language Models (GISP). *arXiv:2510.18030*.
- Xia, M., Zhong, Z., & Chen, D. (2022). Structured Pruning Learns Compact and Accurate Models (CoFi). *ACL*.